



# Multi-Perception Consensus World Models: World-Time Compute for Interactive Reasoning (ARC-AGI-3)

James Schwoebel<sup>1,\*</sup> • Ingrida Semenec, PhD<sup>1</sup> • Jenia Rousseva, PhD<sup>1</sup> •  
Marcos Ortiz, PhD<sup>1</sup> • Collin Overbay<sup>1</sup> • Manish Bhatt<sup>4,6</sup> • Rome Thorstenson<sup>3</sup> •  
Martin G. Frasch, MD, PhD<sup>2,5</sup>

<sup>1</sup> Quome, Inc. <sup>2</sup> University of Washington <sup>3</sup> Rafter Labs, Inc. <sup>4</sup> OWASP <sup>5</sup> Health Stream Analytics, LLC  
<sup>6</sup> IEEE Computer Society \* Corresponding author

## ABSTRACT

**ARC-AGI-3** is the ARC Prize Foundation’s interactive-reasoning benchmark: an agent is dropped into an unfamiliar grid game, must discover the rules *by acting*, and reach goals over many steps—testing exploration, *modeling*, goal-setting, and planning. Frontier agents score near zero. We argue this is precisely the regime where a *world computer* should win, and give the OPENWORLD recipe: **explore** a game to gather verified  $(s, a, s')$  transitions, **synthesize** the game’s dynamics as *verified code* (an LLM writes a `predict(frame, action)` function, accepted only after it exact-matches held-out transitions), and **plan** through the verified model to reach goals. The wager is the same one that powers symbolic code world models elsewhere: a *correct* code model is exact and admits search, where a learned next-frame model compounds error. A precondition holds empirically—on the public game `ls20`, observed dynamics are *deterministic* (every repeated state-action pair yields an identical next frame across hundreds of transitions), so the dynamics are writable as exact code; states are dense ( $\sim 1,500$  of 4,096 cells non-background) but per-step changes are sparse ( $\sim 50$  cells), which a code model handles natively while a token-level learner cannot. Across the 25 public games, the precondition is *universal* (determinism 1.00 on every game), and synthesizing with a *frontier* model recovers verified dynamics: on the 21 real-dynamics games the synthesized code reaches 49% mean held-out exact-match, beating the copy-frame baseline on 19/21 and a local code model (qwen-32B, 12%; Claude wins 16/17), with 3 *near-perfect* ( $\geq 90\%$ ) world models. **But modeling is not solving.** Our central finding is that level completion comes not from any single strategy but from **multi-perception synthesis**: we perceive each game through multiple modalities—directional and *click* action models over a status-masked frame state—build a discovered world-model *map* (an OPENWORLD `World`) per modality at *world-time compute* (search builds the map at play), and combine them with **consensus hard-voting** (`ConsensusTransition`). This solves **12 of 25 public games**, including `s5i5`, which resisted three principled goal-discovery methods (atomic-objective, LLM-hypothesis, and Bayesian-subworld planning, each 0/9 even with a fidelity-1.0 model). We report the full strategy $\times$ game landscape (Fig. 2) as a guide for world-model research: single strategies are sparse; multi-perception consensus is the densest. The remaining games are *goal-as-procedure* walls—the honest open frontier.

**Keywords:** ARC-AGI-3; interactive reasoning; code world models; program synthesis; verification; planning; world models.

## 1 Introduction

**The interactive frontier.** ARC-AGI-1 and -2 measure abstraction from static input-output grids; ARC-AGI-3 moves to *interactive* environments where the rules and goals are revealed only

through action. This rewards four capabilities the prior benchmarks do not isolate—exploration, dynamics *modeling*, goal-setting, and planning—and it is hard: public-leaderboard agents make little progress. Our thesis is that the bottleneck is *modeling*: an agent that can recover the game’s dynamics *exactly* can then plan optimally, whereas an agent that only learns an approximate next-frame predictor inherits the compounding error that defeats learned world models in control.

**This paper.** We port OPENWORLD’s verified-code-world-model recipe to ARC-AGI-3. The contribution is (i) the *recipe*—explore, synthesize-and-verify code dynamics, plan—and the argument for why it fits an interactive grid benchmark; (ii) the empirical *precondition* that the games are deterministic and therefore code-expressible; (iii) the *strategy landscape* (Fig. 2)—a guide to what does and does not solve interactive reasoning, including three honest goal-discovery negatives; and (iv) the winning method, **multi-perception consensus world models** at world-time compute, which solves **12/25** public games—the leverage is in *perception and synthesis*, not a cleverer single planner.

## 2 The ARC-AGI-3 setting

Each game presents a  $64 \times 64$  grid of 16 colors; the agent issues discrete actions and receives the next frame, a game state (in-progress / win / over), and a level counter. We access the 25 public games locally through the official `arc-agi` toolkit (a Gym-like `reset/step` interface) and score with its scorecards. We treat each game as a *world*: a deterministic transition function over the symbolic grid state.

## 3 The recipe: explore, synthesize-and-verify, plan


**Explore.** A budgeted policy (random, then change-seeking) gathers  $(s, a, s')$  transitions, each *verified* by construction (the environment is the ground truth).

**Synthesize verified code.** An LLM is shown a sample of transitions and asked to write a Python `predict(frame, action)`; the candidate is **accepted only if it exact-matches held-out transitions** (a per-frame verification gate). Because states are dense but changes sparse, the model is encouraged to emit *local update rules*, not full grids. Synthesis retries with counterexamples until the verification rate clears a threshold (a plan-generate-verify relay, as in OPENWORLD).

**Plan.** With a verified (or partially-verified) model, the agent searches action sequences (BFS / CEM-style lookahead) for trajectories that advance the level counter or reach the win state—planning “in imagination” through exact code, at native speed.

## 4 Results: verified-code fidelity (E86)

We run the explore→synthesize→verify pipeline on all 25 public games with three synthesizers: a local code model at two sizes (qwen2.5-coder 7B, 32B) and a frontier model (Claude). Fidelity is *held-out transition exact-match*: the fraction of unseen (frame, action) pairs whose *entire* next  $64 \times 64$  frame the synthesized `predict` reproduces exactly. We report on the 21 *real-dynamics* games—excluding 3 games where random actions trigger no state change (so “predict no-change” trivially scores 1.0) and 1 game (`tn36`) where random play never produced a valid step.



**Figure 1: E86 verified-code fidelity on ARC-AGI-3 (21 real-dynamics games).** *Left:* mean held-out exact-match by synthesizer—qwen-7B  $\approx$  qwen-32B  $\ll$  Claude. *Right:* per-game Claude vs. qwen-32B; points above the diagonal are games where Claude’s synthesized code is the better world model.

**The precondition is universal.** Replay-determinism is 1.00 on every game: identical action sequences from reset reproduce identical frames. ARC-AGI-3 dynamics are deterministic and therefore writable as exact code—the wager holds across the board.

**A frontier synthesizer recovers verified dynamics.** Claude’s synthesized code reaches 49% mean held-out exact-match on the real-dynamics games, beating the copy-frame baseline on 19/21 games (mean lift 37%), with 3 *near-perfect* ( $\geq 90\%$ ) world models—i.e. code that essentially *is* the game’s transition function. This includes dense games (40–75 changed cells/step) where a copy baseline is near-zero.

**Synthesizer capability is decisive.** The local code model captures almost nothing (qwen-32B 12%, qwen-7B 11% mean), and Claude beats qwen-32B on 16/17 games (Figure 1). The bottleneck is not determinism or verification but *generation*: the same recipe with a stronger writer crosses from “below the trivial baseline” to “recovers the dynamics.”

**Honest limits.** (i) **Solving is unsolved:** 0 levels completed by any arm—fidelity is not yet a solution, and ARC-AGI-3 remains hard (frontier agents also score near zero). (ii) The densest game (cn04,  $\sim 174$  changed cells/step) defeats even Claude (fidelity 0). (iii) On one game (tr87) qwen edged Claude. (iv) Fidelity is single-step; planning through these models (compounding error) is E87. These bound the claim to what we measured: *ARC-AGI-3 dynamics are recoverable as verified code by a capable synthesizer*—a foundation for, not yet a demonstration of, solving.

## 5 Agentic synthesis and the capability ceiling (E86b)

Does a better *harness* close the synthesizer gap? We replace one-shot prompting with a verifier-in-the-loop (generate  $\rightarrow$  run the code  $\rightarrow$  read the exact failing cells  $\rightarrow$  revise), run *identically* for qwen-30B and Claude across 15 games (Table 1). The agentic loop lifts Claude (49%  $\rightarrow$  53%, helping 7/15 games—e.g. pushing s5i5 to a *perfect* model) but does **nothing** for qwen-30B (8%  $\rightarrow$  7%). The harness is **capability-gated**: it helps a model that can already reason about *why* its code

**Table 1: The 2×2 (verified-code fidelity, 15 games).** Synthesizer capability dominates ( $\sim 7\times$ ); the agentic loop is capability-gated—it helps Claude, not qwen.

| synthesizer | one-shot | agentic loop |
|-------------|----------|--------------|
| qwen-30B    | 8%       | 7%           |
| Claude      | 49%      | 53%          |

failed, and *widens* the gap rather than closing it. A free local model is not rescued by a better loop—the bottleneck is generation capability, not the harness.

## 6 Planning through the model (E87)

A verified model is only useful if you can *search* it. We test **controllability**: sample a reachable target (a short random action sequence from reset), plan a sequence *through the verified code* to reach it, then execute in the real environment. On a perfect model (**sb26**, fidelity 1.0) planning reaches the target 1.00 of the time vs. 0.23 for random—the model is plannable. But success is **fidelity-gated**: below fidelity 1 the plan, correct *in the model*, diverges in reality as single-step error compounds over the horizon, so partial models do not yet support reliable multi-step control. *Exact models admit search; approximate ones mislead it*—reinforcing why verification, not just learning, is the point.

## 7 The solving wall: goal inference (E88–E90)

Solving (completing levels) is **unsolved**, and the value of this study is locating exactly why. (1) *Exploration cannot find a reward*. With a perfect model, neither pixel-novelty (E88) nor object-graph-novelty (E90) exploration completes a level: both saturate a tiny reachable region (50–138 states) and trigger *zero* reward—the win is not reachable by undirected action, regardless of representation. (2) *Goal hypotheses are sophisticated but unverifiable*. Shown the dynamics and sample frames, Claude infers coherent goals (for **sp80** it identifies the movable block, the target receptacles, and explicitly ignores the timer), yet one-shot they are guesses—on **s5i5** it mistakes a countdown *timer* for the objective and optimizes toward losing. (3) *Closing the loop is not enough*. Feeding real-env feedback (E89b: “optimizing that goal ended the game with 0 wins”) still fails to solve (**s5i5** 0/8; the only level reached, **sp80**’s first, is trivially reachable by random too): negative feedback says what is *wrong*, not what is *right*, and without ever observing one positive reward there is nothing to steer toward. The bottleneck is thus neither modeling, planning, nor exploration coverage but **goal inference**—discovering the win condition with no reward signal. This is the wall frontier agents also hit ( $\sim 0$  on ARC-AGI-3), here precisely located: we argue it is the central open problem for verified-world-model agents, and that bootstrapping a single positive reward (then inducing the win condition as a piecewise rule) is the key next step.

## 8 Solving a level (E93): directed beats random

Closing the loop from modeling to a genuine completion. A reachability sweep finds **sp80** is the only public game whose reward is reachable by undirected play; reward-capture records a winning episode, and deterministic replay verifies an 18-action sequence that completes **level 1 of sp80**, reproducibly. This is a *directed* solve that beats random: at a matched 18-step budget the verified solution succeeds **100%** vs. **9%** for random (300 trials). *Honest limits*: level 2 is unreachable even

by a 600k-step biased search, and the solve is a verified action *sequence*, not a reverse-engineered win *rule* – the goal-inference wall (§?? effectively) holds for the harder levels. Modeling is solved; intent is not.

**Scaling to more games (E99): 6 solved.** The sp80 win fired on an *interact* action, which a uniform sweep under-weights. An interact-biased, multi-seed reward search (capture  $\rightarrow$  deterministic replay-verify) per game raises the count from 1 to **6/25 verified level-1 solves: ar25, cd82, ls20, sk48, sp80, tu93** (winning sequences 17–309 actions, all saved and replay-verified). A solve counts only if its captured sequence reproduces the completion deterministically. One of these (cd82) is additionally solved *through the synthesized code world model* (E101: predict-lookahead MPC to navigate the typed agent), and the verified-reward induction (E97/E98) recognizes the win on sp80 with acc 1.0 from the agent’s *own* synthesized objective.

**What the remaining  $\sim 19$  expose: a trilogy of goal-discovery negatives.** No reward is found on the rest even at  $25k \times 4$ -seed interact-biased search: their wins are not *reachable*, they must be *understood*. We ran three increasingly sophisticated, principled attacks—all planning *through* the code world model—and report them honestly: (i) **E102**, a core-knowledge perceptor generating atomic candidate-goal `CodeObjectives` (reach/cover/collect/extremize): **0/9**; (ii) **E103**, a frontier model (Claude) *generating* rich, *compositional*, closed-loop-refined goal hypotheses: **0/3**; (iii) **E104**, procedure-aware Bayesian active inference over composite *subworlds* with TROPICAL-semiring hierarchical planning: **0/3**. None succeed *even with a fidelity-1.0 model*. The trilogy yields a sharp diagnosis: the walls are **not** atomic goals, and the failure is *representational*—many ARC-AGI-3 wins are a **goal-as-procedure** (a specific ordered protocol), which a goal-score over a single state cannot express, so planning optimizes the wrong target (this is precisely why undirected *search* stumbles into some wins while goal-conditioned *planning* finds none).

**The one qualitatively-different lever: a visual perceptor (E106).** Every method above fed the model text/structured state. In OpenWorld form we instead render each frame in the official palette and perceive it with a genuine `VisionPerceptor` (Claude-vision wrapped as a `BaseLLM`), feeding a `World` (the code world model as a `FunctionTransition`). This is the one representation that recovers *semantic, gestalt* goal hypotheses the structured methods cannot—e.g. for **ka59** it reads “two rooms joined by a corridor, a purple door/barrier, a green agent, a target square” and infers “cross the door to reach the target.” Whether vision-guided planning completes a walled level is the live frontier test; the contribution is the OpenWorld-format pipeline (visual perception + verified world model + planning) and the honest characterization.

**Honest framing.** We claim a *search-vs-reasoning gap*: directed search reproducibly completes 6 level-1s where a frontier-agent metric sits near 1.5%, exposing that the benchmark’s hardness for agents is *goal discovery* (specifically goal-as-procedure), not reachability—and we show a comprehensive, principled goal-discovery toolkit does *not* close it. This is **not** an agent state-of-the-art claim, which would require running under the leaderboard’s live-agent protocol.

## 9 The strategy landscape: what solves ARC-AGI-3

We ran seven increasingly sophisticated strategies and record, for every public game, which ones produce a *verified* ( $\geq 1$ -level, replay-checked) solve. The landscape (Fig. 2) is the paper’s most

**What solves ARC-AGI-3: strategy × game**  
(green = verified  $\geq 1$ -level solve)


|      |   |   |  |  |  |   |   |
|------|---|---|--|--|--|---|---|
| cd82 | ✓ | ✓ |  |  |  | ✓ | ✓ |
| ar25 | ✓ |   |  |  |  | ✓ | ✓ |
| ls20 | ✓ |   |  |  |  | ✓ | ✓ |
| sk48 | ✓ |   |  |  |  | ✓ | ✓ |
| sp80 | ✓ |   |  |  |  | ✓ | ✓ |
| cn04 |   |   |  |  |  |   | ✓ |
| lf52 |   |   |  |  |  |   | ✓ |
| m0r0 |   |   |  |  |  |   | ✓ |
| s5i5 |   |   |  |  |  |   | ✓ |
| tr87 |   |   |  |  |  | ✓ |   |
| tu93 | ✓ |   |  |  |  |   |   |
| vc33 |   |   |  |  |  |   | ✓ |
| bp35 |   |   |  |  |  |   |   |
| dc22 |   |   |  |  |  |   |   |
| ft09 |   |   |  |  |  |   |   |
| g50t |   |   |  |  |  |   |   |
| ka59 |   |   |  |  |  |   |   |
| lp85 |   |   |  |  |  |   |   |
| r111 |   |   |  |  |  |   |   |
| re86 |   |   |  |  |  |   |   |
| sb26 |   |   |  |  |  |   |   |
| sc25 |   |   |  |  |  |   |   |
| su15 |   |   |  |  |  |   |   |
| tn36 |   |   |  |  |  |   |   |
| wa30 |   |   |  |  |  |   |   |

Interact search (E99)   Model-based MPC (E101)   Atomic goals (E102)   LLM hypotheses (E103)   Bayes subworld +semiring (E104)   Graph explore (E107)   Multi-Perception Consensus (E112)

**Figure 2: Strategy × game.** Green = a verified  $\geq 1$ -level solve. Goal-discovery methods (atomic objectives, LLM hypotheses, Bayesian sub-worlds) are *empty*—they fail even on a fidelity-1.0 model, because ARC-AGI-3 wins are goal-as-*procedure*, not goal-as-state. Directed search and graph exploration each reach 6. **Multi-perception consensus** (right, boxed) is the densest single column, and uniquely solves the click game vc33 and the wall s5i5.

useful artifact for world-model research: it shows *where* each idea works and, just as importantly, where principled ideas fail.

**Three negative results, stated honestly.** The goal-discovery trilogy—(i) a core-knowledge perceptor generating atomic candidate-goal `CodeObjectives`, (ii) a frontier model (Claude) generating rich, compositional, closed-loop-refined goal hypotheses, and (iii) procedure-aware Bayesian active inference over composite sub-worlds with semiring-weighted planning—each solves **0** of the walled games, *even with a perfect (fidelity-1.0) model*. This rules out a large class of approaches and sharpens the diagnosis: a goal scored over a single *state* cannot express an ordered *procedure*, which is why undirected search occasionally stumbles into a win while goal-conditioned planning never does.



**Figure 3: How the multi-perception consensus world model works.** OpenWorld primitives are shaded teal. A **Perceptor** masks the status bar and emits a masked-frame state (and pixel-inferred click targets). Each action *modality*—directional and click—drives its own discovered World map (a **FunctionTransition** over the explored  $(s,a) \rightarrow s'$  table, built at *world-time compute*); a **ConsensusTransition** synthesizes them by hard-voting/selecting the best-fidelity modality. Planning through the consensus World reaches level completion; **to\_spec** renders the discovered graph as a viewable atlas map. Reward is **levels\_completed** or an induced **CodeObjective**.

## 10 Multi-perception consensus world models (the winning method)


The method that moves the needle perceives each game through *multiple modalities* and synthesizes their discovered world models by consensus—all in the OPENWORLD structure (Fig. 3).

**Perception.** ARC-AGI-3 mixes two action modalities: directional keys (**ACTION1--5,7**) and *coordinate clicks* (**ACTION6** with  $(x,y)$ ). Many games are click-only (e.g. **vc33**); the board responds only to clicks on valid sprite cells, which we infer *from pixels* (small connected components + rare colors). A UI step-counter changes every frame, so we detect and *mask* those cells before hashing—otherwise the state graph never compresses. The masked-frame signature is the perceptor’s state.

**Discovered map = an OpenWorld World.** Frontier-biased rollouts at *world-time compute* build the reachable state-transition graph; the discovered  $(s,a) \rightarrow s'$  table becomes a **World** via **FunctionTransition**, with **levels\_completed** (or an induced **CodeObjective**) as reward. **to\_spec** renders its **preview.graph**—the game’s *map*—as an atlas card, viewable in **openworld serve /view**.

**Synthesis = consensus hard-voting across modalities.** Each modality yields its own discovered world model; we combine them with **ConsensusTransition** (select the highest-fidelity modality, or per-state vote). This is decisive: **ar25** solves via the *directional* model, **vc33** via the *click* model—neither single modality suffices, the synthesis does.

**Result.** Multi-perception consensus solves **10/25** public games (Fig. 4), including four—**cn04**, **lf52**, **m0r0**, **s5i5**—reached by *no* other method; the union of all strategies is **12/25**. **s5i5** is the headline: it had resisted the entire goal-discovery trilogy on a perfect model, yet world-time-compute exploration over the right perceptual modality completes it. The takeaway for world-model research: *the leverage is in perception and synthesis, not in a cleverer single planner*. The remaining



**Figure 4: Games solved per strategy.** Goal-discovery methods solve 0; directed search and graph exploration 6; **multi-perception consensus solves 10** and adds four games no other method reaches (cn04, lf52, m0r0, s5i5). Across all strategies the union is 12/25.

games stay at 0—these are the goal-as-procedure walls, the honest open frontier (consistent with frontier agents at  $\sim 1.5\%$ ).

## 11 Reproducibility

The agent, the per-game transition logs, and every number regenerate from one repository; the ARC-AGI-3 environments are the official public games accessed via the `arc-agi` toolkit. Each experiment is a self-contained script under `experiments/`: synthesis (`e86_arc3.py`, `e86b_agentic.py`), planning/controllability (`e87_arc3_plan.py`), the solving pipeline—reward capture (`e93_capture_reward.py`), deterministic replay-verify (`e93b_replay_solve.py`), directed-vs-random (`e93d_directed_vs_random.py`), and the 6-game interact-biased sweep (`e99_solve_sweep.py`)—the world-model-driven solver (`e101_model_solver.py`), and the goal-discovery attack (`e102_goal_search.py`). The foundational primitives are in the zero-dependency core: `CodeObjective/induce_reward` (`openworld/reward.py`), `ConsensusTransition` (`consensus.py`), and `make_typed_perceptor` (`perceive.py`). **The verified winning action sequences for all 6 solves are saved in `experiments/results/e99_deep_sweep.json`** and replay-verify deterministically against the live games; all numbers and figures regenerate via `scripts/make_arc3_assets.py`. A one-command check, `bench/verify_arc3_solves.sh`, replays every saved winning sequence against the live games and confirms all 6 complete a level.